

Homework10

Jun Rao

10/18/2020

Hidden Markov Models

The goal of this homework is to fit hidden Markov Models on a real data set and compare them to optimal changepoint models.


```
library(data.table)
library(ggplot2)
library(stats4)
library(changepoint)
```

1. Plot the profile.id=79 and chromosome=2 from the neuroblastoma data set (x=position, y=logratio). How many observations are there?

```
data(neuroblastoma, package="neuroblastoma")
pro.dt <- data.table(neuroblastoma$profiles)[profile.id=="79" & chromosome=="2"]
pro.dt

##      profile.id chromosome  position      logratio
## 1:          79           2    18094  0.200378798
## 2:          79           2   2063049  0.505890930
## 3:          79           2   3098882  0.345964030
## 4:          79           2   7177474  0.449957484
## 5:          79           2   8107742  0.474046599
##  --- 
## 192:         79           2 239227603 -0.045431429
## 193:         79           2 239471307 -0.020340448
## 194:         79           2 240618997 -0.081613766
## 195:         79           2 242024751 -0.007231569
## 196:         79           2 242801018  0.048236186

(gg <- ggplot()+
  geom_point(aes(x=position, y=logratio), data=pro.dt))
```



```
str(pro.dt)

## Classes 'data.table' and 'data.frame': 196 obs. of 4 variables:
## $ profile.id: Factor w/ 575 levels "1","2","4","5",...: 76 76 76 76 76 76 ...
## $ chromosome: Factor w/ 24 levels "1","2","3","4",...: 2 2 2 2 2 2 2 2 2 2 ...
## ...
## $ position : int 18094 2063049 3098882 7177474 8107742 8179390 8916469
9470069 11607188 11885382 ...
## $ logratio : num 0.2 0.506 0.346 0.45 0.474 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

We can find that there are total 196 observations.

2. Use depmixS4 R package to fit a sequence of Gaussian Hidden Markov Models, from nstates=1 to 10. Use the logLik function to compute the negative log likelihood for each model size. Then plot y=negative log likelihood as a function of x=number of states. After what model size is there a kink in the curve? (that is the model size you should select)

```
total_states = 10
loss.dt.list <- list()
for (nstates in 1:total_states) {
  hmm <- depmixS4::depmix(logratio ~ 1, pro.dt, nstates)
  hmm.fit <- depmixS4::fit(hmm)
  loss.dt.list[[paste(nstates)]] <- data.table(nstates, loss=-depmixS4::logLik(hmm.fit))
```

```


}

## converged at iteration 1 with logLik: -165.3519
## converged at iteration 9 with logLik: 188.4415
## converged at iteration 78 with logLik: 200.6406
## converged at iteration 78 with logLik: 270.1472
## converged at iteration 187 with logLik: 280.6051

loss.list <- do.call(rbind, loss.dt.list)

ggplot()+
  geom_line(aes(nstates, loss), data=loss.list) +
  labs(x = "the number of model size", y =" negative_log_likelihood")

```



After model size equal 3, there is a kink in the curve

For the selected model size, plot the segmentation on top of the data (for simplicity use x=index instead of position). Make sure to show the mean (geom_segment), the standard deviation (geom_rect), and the changepoints (geom_vline). How many changepoints are there?

```

hmm <- depmixS4::depmix(logratio ~ 1, pro.dt, 3)
hmm.fit <- depmixS4::fit(hmm)

## converged at iteration 9 with logLik: 251.3187

```

```

resp.dt.list <- list()
for(resp.i in seq_along(hmm.fit@response)){
  param.list <- hmm.fit@response[[resp.i]][[1]]@parameters
  resp.dt.list[[paste(resp.i)]] <- data.table(
    state=resp.i,
    mean=param.list[["coefficients"]],
    sd=param.list[["sd"]])
}
}

resp.dt <- do.call(rbind, resp.dt.list)
post.dt <- data.table(hmm.fit@posterior)

## rle = run Length encoding, for more compressed version of hidden
## markov model, in terms of segments, like optimal changepoint
## models.
state.rle <- rle(post.dt$state)

segs.dt <- data.table(
  n.data=state.rle[["lengths"]],
  state=state.rle[["values"]])
segs.dt[, end := cumsum(n.data)]
segs.dt[, start := c(1, end[-.N]+1)]
segs.with.params <- resp.dt[segs.dt, on="state"]

pro.dt[, index := 1:.N]

ggplot()+
  geom_point(aes(
    index, logratio),
    data=pro.dt)+  


  geom_vline(aes(#changes
    xintercept=end+0.5
  ),
  color="red",
  data=segs.with.params[-.N])+  

  geom_rect(aes(#sd
    xmin=start, xmax=end,
    ymin=mean-sd, ymax=mean+sd
  ),
  data=segs.with.params,
  fill="green",
  alpha=0.2,
)+  

  geom_segment(aes(#mean
    x=start, xend=end,
    y=mean, yend=mean),
  color="green",

```

```
size=2,
data=segs.with.params)
```



We total have 8 segments, and 7 changepoints

4.Run `changepoint::cpt.meanvar(method="SegNeigh", penalty="Manual", Q=the same number of segments as in the HMM)` to compute the corresponding optimal changepoint model, and plot that model on top of the data (same as in problem 3). Which model seems to be a better fit, or are they about the same?

```
n.segs <- 8
segs.dt.list <- list()
optimal.models<-changepoint::cpt.meanvar(pro.dt$logratio,penalty="Manual",method="SegNeigh",Q=n.segs)

## Warning in changepoint::cpt.meanvar(pro.dt$logratio, penalty = "Manual", :
## SegNeigh is computationally slow, use PELT instead

## Warning in segneigh.meanvar.norm(c(0.200378797984026, 0.505890929729957, :
The
## number of segments identified is Q, it is advised to increase Q to make sure
## changepoints have not been missed.

end <- cpts(optimal.models)
end <- append(end, length(pro.dt$logratio), after = length(end))
```


```

start <- c(1, end[-length(end)]+1)

seg$dt.list[[paste(n.segs)]] <- data.table(start, end)[, .(
  segments=n.segs,
  mean=mean(pro.dt$logratio[start:end]),
  sd = sd(pro.dt$logratio[start:end])
), by=.(start, end)]
seg$dt <- do.call(rbind, seg$dt.list)

ggplot()+
  geom_point(aes(
    x=index, y=logratio),
    data=pro.dt)+
  geom_segment(aes(
    x=start, y=mean,
    xend=end, yend=mean),
    color="green",
    size = 2,
    data=seg$dt)+
  geom_vline(aes(
    xintercept=start-0.5),
    color="red",
    data=seg$dt[start>1])

```



There are almost the same. There is a liitle bit difference.But I thought HMM algorithm fits it better.

The goal is to systematically compare the negative log likelihood of the two packages (changepoint and depmixS4). changepoint::cpt.meanvar with Q segments should always compute the optimal sequence of Q segment mean/variance parameters (Q-1 changepoints), using dynamic programming. depmixS4/HMM always has nstates <= number of segments, so when there are more segments than states the dynamic programming should have a larger likelihood (better fit to data).

Run HMM for nstates=1 to 10, and compute the number of segments in each model. For each model use changepoint::cpt.meanvar(Q=that number of segments).

Save segment mean/sd/start/end parameters for both algorithms and all model sizes in a single data table.

```
result.dt.list <- list()

total_states <- 9
for (nstates in 2:total_states){

  opt.segs.dt.list <- list()

  hmm <- depmixS4::depmix(logratio ~ 1, pro.dt, nstates)
  hmm.fit <- depmixS4::fit(hmm)

  resp.dt.list <- list()
  for(resp.i in seq_along(hmm.fit@response)){
    param.list <- hmm.fit@response[[resp.i]][[1]]@parameters
    resp.dt.list[[paste(resp.i)]] <- data.table(
      state=resp.i,
      mean=param.list[["coefficients"]],
      sd=param.list[["sd"]])
  }

  resp.dt <- do.call(rbind, resp.dt.list)
  post.dt <- data.table(hmm.fit@posterior)

  state.rle <- rle(post.dt$state)

  segs.dt <- data.table(
    n.data=state.rle[["lengths"]],
    state=state.rle[["values"]],
    algorithm = 'HMM')
  segs.dt[, end := cumsum(n.data)]
  segs.dt[, start := c(1, end[-N]+1)]
  segs.with.params <- resp.dt[segs.dt, on="state"]
```

```

pro.dt[, index := 1:N]

hmm.segs.info=segs.with.params[-N]

Nsegments <- length(hmm.segs.info$start)

for (seg.index in 1:Nsegments) {
  start <- hmm.segs.info$start[seg.index]
  end <- hmm.segs.info$end[seg.index]
  mean <- hmm.segs.info$mean[seg.index]
  sd <- hmm.segs.info$sd[seg.index]
  algorithm <- hmm.segs.info$algorithm[seg.index]
  result.dt.list[[paste(nstates, seg.index, start, end, mean, sd, algorithm)]]
<- data.table(nstates, seg.index, start, end, mean, sd, algorithm)
}

optimal.models<-changepoint::cpt.meanvar(pro.dt$logratio, penalty="Manual",
method="SegNeigh", Q=Nsegments)
opt.end <- cpts(optimal.models)
opt.end <- append(opt.end, length(pro.dt$logratio), after = length(opt.end))

opt.start <- c(1, opt.end[-length(opt.end)]+1)

opt.segs.dt.list[[paste(Nsegments)]] <- data.table(opt.start, opt.end)[,
.(
  segments=Nsegments,
  mean=mean(pro.dt$logratio[opt.start:opt.end]),
  sd = sd(pro.dt$logratio[opt.start:opt.end]),
  algorithm = 'OPT'
), by=(opt.start, opt.end)]
opt.segs.info <- do.call(rbind, opt.segs.dt.list)

for (seg.index in 1:Nsegments) {
  start <- opt.segs.info$opt.start[seg.index]
  end <- opt.segs.info$opt.end[seg.index]
  mean <- opt.segs.info$mean[seg.index]
  sd <- opt.segs.info$sd[seg.index]
  algorithm <- opt.segs.info$algorithm[seg.index]
  result.dt.list[[paste(nstates, seg.index, start, end, mean, sd, algorithm)]]
<- data.table(nstates, seg.index, start, end, mean, sd, algorithm)
}
}

head(result.dt <- do.call(rbind, result.dt.list))

```

	nstates	seg.index	start	end	mean	sd	algorithm
## 1:	2	1	1	33	0.755056018	1.13626391	HMM
## 2:	2	2	34	135	0.009681706	0.04711407	HMM
## 3:	2	3	136	136	0.755056018	1.13626391	HMM
## 4:	2	4	137	154	0.009681706	0.04711407	HMM
## 5:	2	5	155	155	0.755056018	1.13626391	HMM
## 6:	2	1	11	11	0.401781188	0.08316934	OPT

Compute the negative log likelihood using stats::dnorm function, and plot it as a function of the number of segments, different algorithms in different colors (e.g., red=HMM, black=DP). Is the curve lower for DP as expected?

```

hmm.result.dt <- subset(result.dt,result.dt$algorithm == 'HMM')
opt.result.dt <- subset(result.dt,result.dt$algorithm == 'OPT')

neg.log.likelihood.list <- list()

for (nstates_index in 2:total_states){

  hmm.mean.dt <- subset(hmm.result.dt,hmm.result.dt$nstates==nstates_index)$mean
  neg_log_L = -sum(dnorm(hmm.mean.dt, mean = mean(hmm.mean.dt), sd = sd(hmm.mean.dt)))
  algorithm <- 'HMM'
  neg.log.likelihood.list[[paste(nstates_index,neg_log_L,algorithm)]] <- data.table(nstates_index,neg_log_L,algorithm)


  opt.mean.dt <- subset(opt.result.dt,opt.result.dt$nstates==nstates_index)$mean
  neg_log_L = -sum(dnorm(opt.mean.dt, mean = mean(opt.mean.dt), sd = sd(opt.mean.dt)))
  algorithm <- 'DP'
  neg.log.likelihood.list[[paste(nstates_index,neg_log_L,algorithm)]] <- data.table(nstates_index,neg_log_L,algorithm)

}

neg.log.likelihood.dt <- do.call(rbind, neg.log.likelihood.list)

ggplot()+
  geom_line(aes(
    nstates_index, neg_log_L, color=algorithm),
    data= neg.log.likelihood.dt)+
  xlab("the number of model size") + ylab("negative log likelihood")

```



These two algorithms are almost same. DP performance as we expected