

Homework11

Jun Rao

10/26/2020

Segmentation model selection and evaluation

The goals of this homework are to explore how a subtrain/validation split can be used to select the segmentation model size parameter. compare segmentation models in terms of label error rates and ROC curves

```
library(ggplot2)
library(data.table)
data(neuroblastoma, package="neuroblastoma")
```

1. For data sequence profile.id=79, chromosome=2 in the neuroblastoma data, assign observations to 50% subtrain 50% validation sets. How many observations are in each set?
Hint: use table(set)

```
pro.dt <- data.table(neuroblastoma$profiles)[profile.id=="79" & chromosome=="2"]
pro.dt[, set := rep(rep(c("subtrain","validation"), each=2), l=.N)]
pro.dt[, index := 1:.N]

subtrain.dt <- pro.dt[set=="subtrain"]
validation.dt <- pro.dt[set=="validation"]

str(subtrain.dt)

## Classes 'data.table' and 'data.frame': 98 obs. of 6 variables:
## $ profile.id: Factor w/ 575 levels "1","2","4","5",...: 76 76 76 76 76 76 ...
## $ chromosome: Factor w/ 24 levels "1","2","3","4",...: 2 2 2 2 2 2 2 2 2 2 ...
...
## $ position : int 18094 2063049 8107742 8179390 11607188 11885382 156479
03 15998838 17554158 18160714 ...
## $ logratio : num 0.2 0.506 0.474 0.43 0.455 ...
## $ set      : chr "subtrain" "subtrain" "subtrain" "subtrain" ...
## $ index    : int 1 2 5 6 9 10 13 14 17 18 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

We total have 196 observations.

In our set, each set have 98 observations, because I use each = 2 at here. It will equally split the original data.

2. Use binsegRcpp::binseg_normal to compute models from 1 to 10 segments on the subtrain data, then compute one data table with predicted changepoint positions

```
subtrain.binseg.models <- binsegRcpp::binseg_normal(subtrain.dt$logratio, 10)
subtrain.binseg.dt <- coef(subtrain.binseg.models)

subtrain.getPos <- function(subtrain.index){
  orig.index <- subtrain.dt$index[ subtrain.index ]
  pro.dt$position[ orig.index ]
}

subtrain.binseg.dt[, start_in_all_data := subtrain.getPos(subtrain.binseg.dt$start)]
subtrain.binseg.dt[, end_in_all_data := subtrain.getPos(subtrain.binseg.dt$end)]
subtrain.binseg.dt[, type := rep("subtrain", l=N)]
head(subtrain.binseg.dt)

##   segments start end      mean start_in_all_data end_in_all_data type
## 1:       1    1  98 0.14558115          18094        240618997 subtrain
## 2:       2    1    8 1.30763500          18094        15998838 subtrain
## 3:       2    9   98 0.04228748         17554158        240618997 subtrain
## 4:       3    1    6 0.40969361          18094        11885382 subtrain
## 5:       3    7    8 4.00145916         15647903        15998838 subtrain
## 6:       3    9   98 0.04228748         17554158        240618997 subtrain
```

and another data table with segment start/end values (in terms of the full data indices/positions, not just the subtrain set).

```
validation.binseg.models <- binsegRcpp::binseg_normal(validation.dt$logratio, 10)
validation.binseg.dt <- coef(validation.binseg.models)

validation.getPos <- function(validation.index){
  orig.index <- validation.dt$index[ validation.index ]
  pro.dt$position[orig.index ]
}

validation.binseg.dt[, start_in_all_data := validation.getPos(validation.binseg.dt$start)]
validation.binseg.dt[, end_in_all_data := validation.getPos(validation.binseg.dt$end)]
```

```

.dt$end)]
validation.binseg.dt[, type := rep("validation", l=N)]
head(validation.binseg.dt)

##    segments start end      mean start_in_all_data end_in_all_data      t
ype
## 1:      1     1 98 0.14231781      3098882 242801018 validation
## 2:      2     1  7 1.36618893      3098882 16084178 validation
## 3:      2     8 98 0.04817387      16498022 242801018 validation
## 4:      3     1  5 0.39228628      3098882 12173642 validation
## 5:      3     6  7 3.80094555      15291502 16084178 validation
## 6:      3     8 98 0.04817387      16498022 242801018 validation

```

Plot these models on top of the data (use points with different color/fill for different sets, black=subtrain, red=validation), and use facet_grid(segments ~ .) to draw each model size in a different panel. After how many segments does the model appear to overfit?

```

total.dt <- rbind(subtrain.binseg.dt, validation.binseg.dt)

for(col.name in c("start_in_all_data", "end_in_all_data")){
  col.value <- total.dt[[col.name]]
  set(total.dt, j=paste0(col.name, ".pos"),
      value=pro.dt$position[col.value])
}

total.dt[, end.before := c(NA, end_in_all_data[-N]), by=(type, segments) ]
change.dt <- data.table(pro.dt, total.dt[1 < start_in_all_data])

## Warning in as.data.table.list(x, keep.rownames = keep.rownames, check.names
## = check.names, : Item 2 has 110 rows but longest item has 196; recycled wi
## th
## remainder.

change.dt[, changepoint := (start_in_all_data+end.before)/2]

gg <- ggplot(aes(x=position, y=logratio), data=pro.dt)+geom_point()


(gg+ facet_grid(segments ~ .)+
  geom_segment(aes(
    x=start, y=mean,
    xend=end, yend=mean,
    color=type)),

```

```

    data=total.dt) +
  geom_vline(aes(
    xintercept=end_in_all_data,
    color=type),
    data=change.dt)+ 
  scale_size_manual(values=c(subtrain=2, validation=1))
)

```



In our result, it looks like after 3 segments, the model appears to overfit

3. Compute the square loss for each model size and set, and store these numbers in a data table with columns segments, set,loss. Plot the loss as a function of the number of segments, with different sets in different colors (black=subtrain, red=validation). Based on the loss values/plot, what is the number of segments you should select?

```

subtrain.loss.dt<- data.table()
validation.loss.dt <- data.table()
subtrain.loss.dt[, loss := subtrain.binseg.models$loss]
subtrain.loss.dt[, segments := subtrain.binseg.models$segments]
subtrain.loss.dt[, type := "subtrain"]
validation.loss.dt[, loss := validation.binseg.models$loss]
validation.loss.dt[, segments := validation.binseg.models$segments]
validation.loss.dt[, type := "validation"]
loss.dt <- rbind(subtrain.loss.dt,validation.loss.dt)

```


```
ggplot() +
```

```

geom_line(aes(
  segments, loss, color=type),
  data=loss.dt)+  

xlab("the number of segments") + ylab("the loss value")

```



Based on the picture, I will choose segments = 3.

4. Compute model selection criteria = loss + penalty*segments, for penalty=2 (AIC) and penalty=log(N.data) (BIC). Save these numbers in a single data table with columns segments, crit.name(AIC/BIC), crit.value(numeric value defined in equation above). Then plot them using different colors for different penalties (e.g., AIC=blue, BIC=orange). Do the two criteria select the same number of segments?

```

binseg.models <- binsegRcpp::binseg_normal(pro.dt$logratio, 10)
AIC.criteria.dt<- data.table()
BIC.criteria.dt <- data.table()

AIC.criteria.dt[, criteria.value := binseg.models$loss + 2 * binseg.models$segments]
AIC.criteria.dt[, segments := binseg.models$segments]
AIC.criteria.dt[, criteria.name := "AIC"]

BIC.criteria.dt[, criteria.value := binseg.models$loss + log(length(pro.dt$logratio)) * binseg.models$segments]
BIC.criteria.dt[, segments := binseg.models$segments]
BIC.criteria.dt[, criteria.name := "BIC"]


```

```

criteria.dt <- rbind(AIC.criteria.dt,BIC.criteria.dt)

ggplot()+
  geom_line(aes(
    segments, criteria.value, color=criteria.name),
    data=criteria.dt)+
  xlab("the number of segments") + ylab("the criteria value")

```



Yes. The two criteria select the same number of segments which is 3.

The goal is to run binsegRcpp::binseg_normal on ALL labeled neuroblastoma data sequences, and compare two model selection criteria (AIC/BIC) in terms of label error rates and ROC curves. You should have a for loop over rows of the “annotations” data table, since there is only one label/row per data sequence. For each sequence (unique value of profile.id, chromosome) you should run binary segmentation up to 20 segments. Use penaltyLearning::modelSelection to compute a data table representing the model selection function (which maps penalty values to model sizes).

```

AIC.loss.dt.list <- list()
BIC.loss.dt.list <- list()
change.dt.list <- list()

all.profiles <- data.table(neuroblastoma$profiles)
all.labels <- data.table(neuroblastoma$annotations)
label.i.todo <- 1:nrow(all.labels)

```

```

# Label.i.done <- as.integer(unique(sub(" .*", "", names(loss.dt.list))))
# Label.i.new <- Label.i.todo[! Label.i.todo %in% Label.i.done]

max.segments <- 10

for (label.i in label.i.todo) {
  # cat(sprintf("Label.i=%d\n", label.i))
  one.label <- all.labels[label.i]
  select.dt <- one.label[, .(profile.id, chromosome)]
  pro.dt <- all.profiles[select.dt, on=names(select.dt)]

  this.max <- if(nrow(pro.dt) < max.segments){
    nrow(pro.dt)
  }else{
    max.segments
  }

  binseg.models <- binsegRcpp::binseg_normal(pro.dt$logratio, this.max)

  segs.dt.list <- list()

  for (n.segs in 1:this.max) {
    end <- binseg.models$end[1:n.segs]
    start <- c(1, end[-length(end)]+1)
    segs.dt.list[[paste(n.segs)]] <- data.table(start, end)[, .(
      segments=n.segs,
      mean=mean(pro.dt$logratio[start:end]),
      algorithm="DP"
    ), by=.(start, end)]
  }

  segs.dt <- do.call(rbind, segs.dt.list)

  for(col.name in c("start", "end")){
    col.value <- segs.dt[[col.name]]
    set(segs.dt, j=paste0(col.name, ".pos"),
        value=pro.dt$position[col.value])
  }

  segs.dt[, end.before := c(NA, end.pos[-N]), by=.(segments) ]
  change.dt <- data.table(select.dt, segs.dt[1 < start])
  change.dt[, changepoint := (start.pos+end.before)/2]
  AIC.this.loss.dt <- data.table(
    segments=1:this.max,
    loss=binseg.models$loss)

  BIC.this.loss.dt <- data.table(

```

```

    segments=1:this.max,
    loss=binseg.models$loss)

penalty <- 0.12
AIC.this.loss.dt[, crit.value := loss + penalty*segments]
BIC.this.loss.dt[, crit.value := loss + log10(length(pro.dt$logratio))*segments]
AIC.this.loss.dt[, criteria.name := "AIC"]
BIC.this.loss.dt[, criteria.name := "BIC"]
AIC.loss.dt.list[[paste(label.i)]] <- data.table(
  select.dt, AIC.this.loss.dt)
change.dt.list[[paste(label.i)]] <- change.dt[, data.table(
  select.dt, changepoint, segments)]

BIC.loss.dt.list[[paste(label.i)]] <- data.table(
  select.dt, BIC.this.loss.dt)
change.dt.list[[paste(label.i)]] <- change.dt[, data.table(
  select.dt, changepoint, segments)]

}

change.dt <- do.call(rbind, change.dt.list)
AIC.loss.dt <- do.call(rbind, AIC.loss.dt.list)
BIC.loss.dt <- do.call(rbind, BIC.loss.dt.list)

## Compute model selection function, which maps penalty (Lambda)
## values to model complexity (segments) values.
AIC.all.model.selection <- AIC.loss.dt[, penaltyLearning::modelSelection(
  .SD, "loss", "segments"),
  by=.(profile.id, chromosome)]
pred.penalty.dt <- AIC.loss.dt[, data.table(
  pred.log.lambda=log(10)
), by=.(profile.id, chromosome)]

## Compute Label error, fp/fn for each selected model.
AIC.error.list <- penaltyLearning::labelError(
  models=AIC.all.model.selection,
  labels=all.labels[label.i.todo],
  changes=change.dt,
  problem.vars=c("profile.id", "chromosome"),
  change.var="changepoint",
  model.vars="segments")
AIC.error.list$label.errors[, .(
  profile.id, chromosome, min.lambda, max.lambda, segments, fp, fn)]

```

```

##      profile.id chromosome min.lambda max.lambda segments fp fn
##      1:           1          1 0.00000000 0.10981858     10  1  0
##      2:           1          1 0.10981858 0.14642178      7  1  0
##      3:           1          1 0.14642178 0.27946950      5  1  0
##      4:           1          1 0.27946950 1.23652855      4  0  0
##      5:           1          1 1.23652855 1.83184494      3  0  0
##      ---
##  20062:       603          17 0.06763515 0.07307963      7  0  0
##  20063:       603          17 0.07307963 0.11831418      6  0  0
##  20064:       603          17 0.11831418 0.17108867      5  0  0
##  20065:       603          17 0.17108867 2.10773718      2  0  1
##  20066:       603          17 2.10773718        Inf      1  0  1

## Compute ROC curve, FPR/TPR for every prediction threshold (default
## prediction threshold is zero).
AIC.roc.list <- penaltyLearning::ROChange(
  AIC.error.list$model.errors,
  pred.penalty.dt,
  problem.vars=c("profile.id", "chromosome"))

AIC.roc.list$roc[, .(min.thresh, max.thresh, FPR, TPR, errors)]
##      min.thresh max.thresh      FPR      TPR errors
##  1:      -Inf   -7.526805 0.9785589 1.000000000  2784
##  2:   -7.526805   -7.475998 0.9782074 1.000000000  2783
##  3:   -7.475998   -7.366749 0.9778559 1.000000000  2782
##  4:   -7.366749   -7.362607 0.9775044 1.000000000  2781
##  5:   -7.362607   -7.337930 0.9771529 1.000000000  2780
##  ---
## 3321:    1.943780   2.099433 0.0000000 0.006980803  569
## 3322:    2.099433   2.121908 0.0000000 0.005235602  570
## 3323:    2.121908   2.338209 0.0000000 0.003490401  571
## 3324:    2.338209   2.491455 0.0000000 0.001745201  572
## 3325:    2.491455        Inf 0.0000000 0.000000000  573

## Compute model selection function, which maps penalty (Lambda)
## values to model complexity (segments) values.
BIC.all.model.selection <- BIC.loss.dt[, penaltyLearning::modelSelection(
  .SD, "loss", "segments"),
  by=.(profile.id, chromosome)]
pred.penalty.dt <- BIC.loss.dt[, data.table(
  pred.log.lambda=log(10)
), by=.(profile.id, chromosome)]

## Compute Label error, fp/fn for each selected model.
BIC.error.list <- penaltyLearning::labelError(
  models=BIC.all.model.selection,
  labels=all.labels[label.i.todo],

```

```

changes=change.dt,
problem.vars=c("profile.id", "chromosome"),
change.var="changepoint",
model.vars="segments")
BIC.error.list$model.errors[, .(
  profile.id, chromosome, min.lambda, max.lambda, segments, fp, fn)]

##      profile.id chromosome min.lambda max.lambda segments fp fn
## 1:          1          1  0.00000000  0.10981858     10  1  0
## 2:          1          1  0.10981858  0.14642178      7  1  0
## 3:          1          1  0.14642178  0.27946950      5  1  0
## 4:          1          1  0.27946950  1.23652855      4  0  0
## 5:          1          1  1.23652855  1.83184494      3  0  0
## ...
## 20062:      603          17  0.06763515  0.07307963      7  0  0
## 20063:      603          17  0.07307963  0.11831418      6  0  0
## 20064:      603          17  0.11831418  0.17108867      5  0  0
## 20065:      603          17  0.17108867  2.10773718      2  0  1
## 20066:      603          17  2.10773718        Inf      1  0  1

## Compute ROC curve, FPR/TPR for every prediction threshold (default
## prediction threshold is zero).
BIC.roc.list <- penaltyLearning::ROChange(
  BIC.error.list$model.errors,
  pred.penalty.dt,
  problem.vars=c("profile.id", "chromosome"))
BIC.roc.list$roc[, .(min.thresh, max.thresh, FPR, TPR, errors)]
```

	min.thresh	max.thresh	FPR	TPR	errors
1:	-Inf	-7.526805	0.9785589	1.000000000	2784
2:	-7.526805	-7.475998	0.9782074	1.000000000	2783
3:	-7.475998	-7.366749	0.9778559	1.000000000	2782
4:	-7.366749	-7.362607	0.9775044	1.000000000	2781
5:	-7.362607	-7.337930	0.9771529	1.000000000	2780

3321:	1.943780	2.099433	0.0000000	0.006980803	569
3322:	2.099433	2.121908	0.0000000	0.005235602	570
3323:	2.121908	2.338209	0.0000000	0.003490401	571
3324:	2.338209	2.491455	0.0000000	0.001745201	572
3325:	2.491455	Inf	0.0000000	0.000000000	573

Then use the model selection data table as input to penaltyLearning::labelError, which computes the number of incorrectly predicted labels as a function of the penalty/lambda. Also compute a data table of predicted penalty values for AIC/BIC with two rows and columns sequenceID,crit.name,pred.log.lambda. Hint: for reference/example see <http://members.cbio.mines-paristech.fr/~thocking/change-tutorial/Supervised.html>

```

head(BIC.error.list$label.errors)

##      profile.id chromosome segments min      max pred.changes min.lambda
## 1:          1          1          1      10      0 1250000000                  4  0.0000000
```

```

## 2:      1      1      7    0 1250000000      3  0.1098186
## 3:      1      1      5    0 1250000000      1  0.1464218
## 4:      1      1      4    0 1250000000      0  0.2794695
## 5:      1      1      3    0 1250000000      0  1.2365286
## 6:      1      1      2    0 1250000000      0  1.8318449
##   max.lambda min.log.lambda max.log.lambda cum.iterations      loss crit.v
alue
## 1: 0.1098186          -Inf     -2.2089256      12 3.434715 30.1
9250
## 2: 0.1464218        -2.2089256     -1.9212639      7 3.764170 22.4
9462
## 3: 0.2794695        -1.9212639     -1.2748621      4 4.057014 17.4
3591
## 4: 1.2365286        -1.2748621      0.2123079      3 4.336483 15.0
3960
## 5: 1.8318449        0.2123079      0.6053236      2 5.573012 13.6
0035
## 6: 8.5101305        0.6053236      2.1412573      1 7.404857 12.7
5641
##   criteria.name annotation min.changes max.changes color possible.fn
## 1:       BIC      normal         0          0 #f6f4bf      0
## 2:       BIC      normal         0          0 #f6f4bf      0
## 3:       BIC      normal         0          0 #f6f4bf      0
## 4:       BIC      normal         0          0 #f6f4bf      0
## 5:       BIC      normal         0          0 #f6f4bf      0
## 6:       BIC      normal         0          0 #f6f4bf      0
##   possible.fp weight fp fn      status
## 1:      1     1  1  0 false positive
## 2:      1     1  1  0 false positive
## 3:      1     1  1  0 false positive
## 4:      1     1  0  0      correct
## 5:      1     1  0  0      correct
## 6:      1     1  0  0      correct

head(AIC.error.list$label.errors)

##   profile.id chromosome segments min      max pred.changes min.lambda
## 1:      1           1       10    0 1250000000      4  0.0000000
## 2:      1           1       7    0 1250000000      3  0.1098186
## 3:      1           1       5    0 1250000000      1  0.1464218
## 4:      1           1       4    0 1250000000      0  0.2794695
## 5:      1           1       3    0 1250000000      0  1.2365286
## 6:      1           1       2    0 1250000000      0  1.8318449
##   max.lambda min.log.lambda max.log.lambda cum.iterations      loss crit.v
alue
## 1: 0.1098186          -Inf     -2.2089256      12 3.434715 4.63
4715
## 2: 0.1464218        -2.2089256     -1.9212639      7 3.764170 4.60
4170
## 3: 0.2794695        -1.9212639     -1.2748621      4 4.057014 4.65

```

```

7014
## 4: 1.2365286      -1.2748621      0.2123079      3 4.336483 4.81
6483
## 5: 1.8318449      0.2123079      0.6053236      2 5.573012 5.93
3012
## 6: 8.5101305      0.6053236      2.1412573      1 7.404857 7.64
4857
##   criteria.name annotation min.changes max.changes   color possible.fn
## 1:          AIC      normal         0         0 #f6f4bf       0
## 2:          AIC      normal         0         0 #f6f4bf       0
## 3:          AIC      normal         0         0 #f6f4bf       0
## 4:          AIC      normal         0         0 #f6f4bf       0
## 5:          AIC      normal         0         0 #f6f4bf       0
## 6:          AIC      normal         0         0 #f6f4bf       0
##   possible.fp weight fp fn      status
## 1:          1     1 1 0 false positive
## 2:          1     1 1 0 false positive
## 3:          1     1 1 0 false positive
## 4:          1     1 0 0      correct
## 5:          1     1 0 0      correct
## 6:          1     1 0 0      correct

```

Also inspect the thresholds element of the ROChange result in order to see the number of label errors per model. Which model has a smaller number of predicted label errors?

```


AIC.list.dt <- data.table()
AIC.list.dt[, FPR := AIC.roc.list$roc$FPR]
AIC.list.dt[, TPR := AIC.roc.list$roc$TPR]
AIC.list.dt[, name := 'AIC']
BIC.list.dt <- data.table()
BIC.list.dt[, FPR := BIC.roc.list$roc$FPR]
BIC.list.dt[, TPR := BIC.roc.list$roc$TPR]
BIC.list.dt[, name := 'BIC']

roc.list.dt <- data.table()

roc.list.dt <- rbind(AIC.list.dt,BIC.list.dt)

ggplot()+
  geom_path(aes(
    FPR, TPR, color=name),
    data=roc.list.dt)

```



Based on the result, BIC model has a larger area under the curve

Based on the result, AIC model has a smaller number of predicted label errors